

Task 5:- Develop A Neural Network That Can Read Handwriting:

Author:- Mitali D Shinde

Level:-Advanced

```
In [1]: import tensorflow as tf
        from tensorflow.keras import datasets, layers, models
        import matplotlib.pyplot as plt
        %matplotlib inline
        from keras.models import model_from_json

UsageError: line magic function `%matplotlib.inline` not found.
```

Download MNIST dataset

The MNIST database (Modified National Institute of Standards and Technology database) is a large database of handwritten digits that is commonly used for training various image processing systems.

```
In [2]: (train_img, train_labels), (test_img, test_labels) = datasets.mnist.load_data()

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11493376/11490434 [=====] - 5s 0us/step
11501568/11490434 [=====] - 5s 0us/step
```

Normalize Images

```
In [3]: train_img, test_img = train_img/255.0, test_img/255.0
        #normalizing the pixel value between 0 and 1

In [4]: len(train_img)

Out[4]: 60000


In [5]: len(test_img)

Out[5]: 10000
```

Visualizing Images


```
In [6]: plt.imshow(train_img[0], cmap = plt.cm.gray_r, interpolation = 'nearest')

Out[6]: <matplotlib.image.AxesImage at 0x206dad66250>
```




```
In [10]: plt.imshow(train_img[60], cmap = plt.cm.gray_r, interpolation = 'nearest')

Out[10]: <matplotlib.image.AxesImage at 0x206db6341c0>
```



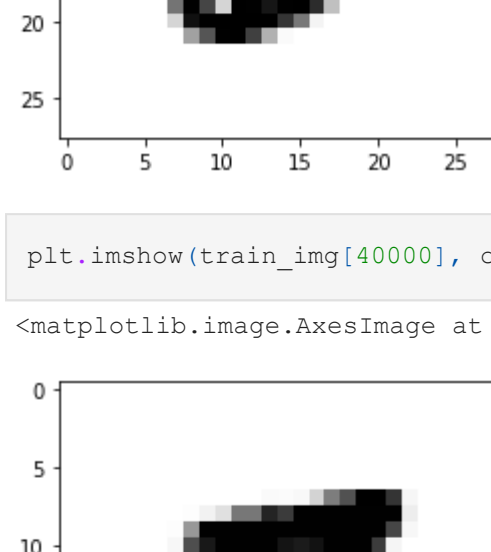
```
In [7]: plt.imshow(train_img[100], cmap = plt.cm.gray_r, interpolation = 'nearest')

Out[7]: <matplotlib.image.AxesImage at 0x206db4f1d00>
```



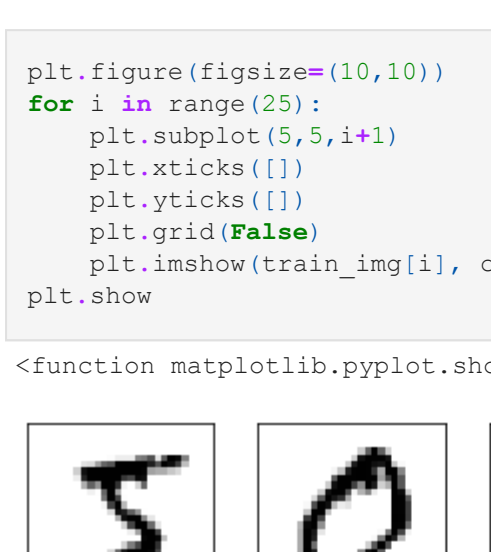
```
In [8]: plt.imshow(train_img[99], cmap = plt.cm.gray_r, interpolation = 'nearest')

Out[8]: <matplotlib.image.AxesImage at 0x206db563f10>
```



```
In [9]: plt.imshow(train_img[999], cmap = plt.cm.gray_r, interpolation = 'nearest')

Out[9]: <matplotlib.image.AxesImage at 0x206db5c3f10>
```



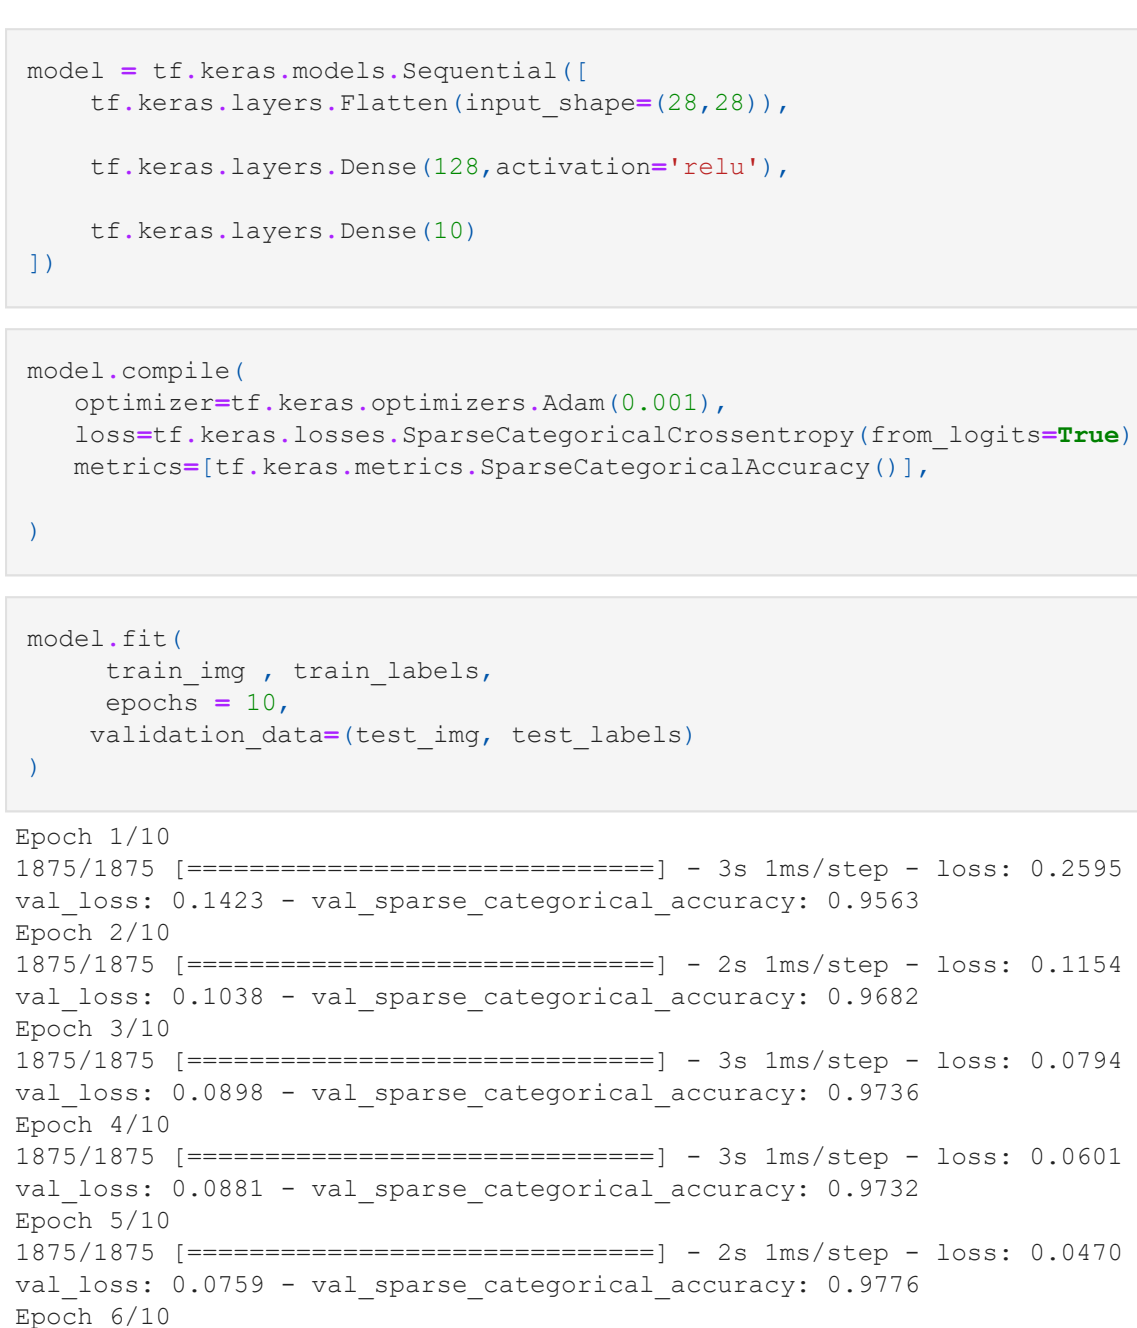
```
In [11]: plt.imshow(train_img[40000], cmap = plt.cm.gray_r, interpolation = 'nearest')

Out[11]: <matplotlib.image.AxesImage at 0x206db6954c0>
```



```
In [12]: plt.figure(figsize=(10,10))
        for i in range(25):
            plt.subplot(5,5,i+1)
            plt.xticks([])
            plt.yticks([])
            plt.grid(False)
            plt.imshow(train_img[i], cmap=plt.cm.binary)
        plt.show

Out[12]: <function matplotlib.pyplot.show(close=None, block=None)>
```



Creating and training the model

Sequential model

```
In [13]: model = tf.keras.models.Sequential([
        tf.keras.layers.Flatten(input_shape=(28,28)),

        tf.keras.layers.Dense(128, activation='relu'),

        tf.keras.layers.Dense(10)
    ])

In [25]: model.compile(
        optimizer=tf.keras.optimizers.Adam(0.001),
        loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
        metrics=[tf.keras.metrics.SparseCategoricalAccuracy()],
    )

In [27]: model.fit(
        train_img, train_labels,
        epochs = 10,
        validation_data=(test_img, test_labels)
    )

Epoch 1/10
1875/1875 [=====] - 3s 1ms/step - loss: 0.2595 - sparse_categorical_accuracy: 0.9258 -
val_loss: 0.1423 - val_sparse_categorical_accuracy: 0.9563
Epoch 2/10
1875/1875 [=====] - 2s 1ms/step - loss: 0.1154 - sparse_categorical_accuracy: 0.9657 -
val_loss: 0.1038 - val_sparse_categorical_accuracy: 0.9682
Epoch 3/10
1875/1875 [=====] - 3s 1ms/step - loss: 0.0794 - sparse_categorical_accuracy: 0.9767 -
val_loss: 0.0898 - val_sparse_categorical_accuracy: 0.9736
Epoch 4/10
1875/1875 [=====] - 3s 1ms/step - loss: 0.0601 - sparse_categorical_accuracy: 0.9812 -
val_loss: 0.0881 - val_sparse_categorical_accuracy: 0.9732
Epoch 5/10
1875/1875 [=====] - 2s 1ms/step - loss: 0.0470 - sparse_categorical_accuracy: 0.9854 -
val_loss: 0.0759 - val_sparse_categorical_accuracy: 0.9776
Epoch 6/10
1875/1875 [=====] - 2s 1ms/step - loss: 0.0365 - sparse_categorical_accuracy: 0.9887 -
val_loss: 0.0820 - val_sparse_categorical_accuracy: 0.9741
Epoch 7/10
1875/1875 [=====] - 3s 1ms/step - loss: 0.0296 - sparse_categorical_accuracy: 0.9906 -
val_loss: 0.0721 - val_sparse_categorical_accuracy: 0.9780
Epoch 8/10
1875/1875 [=====] - 3s 2ms/step - loss: 0.0230 - sparse_categorical_accuracy: 0.9926 -
val_loss: 0.0662 - val_sparse_categorical_accuracy: 0.9804
Epoch 9/10
1875/1875 [=====] - 3s 2ms/step - loss: 0.0195 - sparse_categorical_accuracy: 0.9940 -
val_loss: 0.0864 - val_sparse_categorical_accuracy: 0.9758
Epoch 10/10
1875/1875 [=====] - 3s 2ms/step - loss: 0.0160 - sparse_categorical_accuracy: 0.9948 -
val_loss: 0.0888 - val_sparse_categorical_accuracy: 0.9750
Out[27]: <keras.callbacks.History at 0x206dc9556d0>
```

Model summary

summary() method will display the architecture of the model

```
In [32]: model.summary()

Model: "sequential"

Layer (type) Output Shape Param #
=====
flatten (Flatten) (None, 784) 0
dense (Dense) (None, 128) 100480
dense_1 (Dense) (None, 10) 1290
flatten_1 (Flatten) (None, 10) 0
dense_2 (Dense) (None, 64) 704
dense_3 (Dense) (None, 10) 650
=====
Total params: 103,124
Trainable params: 103,124
Non-trainable params: 0
```

Flatten

```
In [29]: model.add(layers.Flatten())
```

Adding dense layer

```
In [30]: model.add(layers.Dense(64, activation='relu'))
        model.add(layers.Dense(10))
```

Compile

```
In [31]: model.compile(optimizer='adam', loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True), metrics=['sparse_categorical_accuracy'])
```

Training the model

```
In [33]: model.fit(train_img, train_labels, epochs=12, validation_data=(test_img, test_labels))

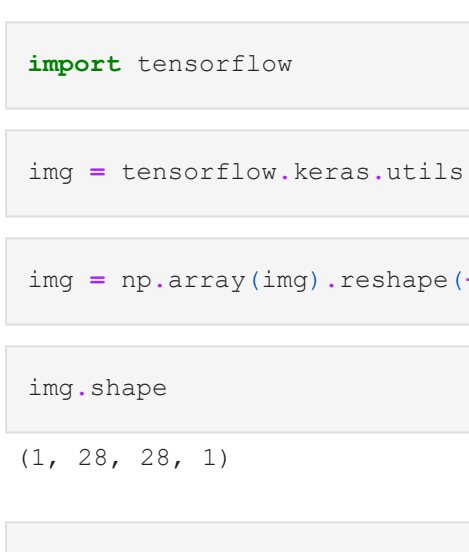
Epoch 1/12
1875/1875 [=====] - 3s 1ms/step - loss: 0.0803 - accuracy: 0.9785 - val_loss: 0.0835 -
val_accuracy: 0.9775
Epoch 2/12
1875/1875 [=====] - 3s 1ms/step - loss: 0.0205 - accuracy: 0.9934 - val_loss: 0.1099 -
val_accuracy: 0.9738
Epoch 3/12
1875/1875 [=====] - 3s 1ms/step - loss: 0.0169 - accuracy: 0.9942 - val_loss: 0.0930 -
val_accuracy: 0.9778
Epoch 4/12
1875/1875 [=====] - 3s 2ms/step - loss: 0.0163 - accuracy: 0.9945 - val_loss: 0.0969 -
val_accuracy: 0.9803
Epoch 5/12
1875/1875 [=====] - 3s 2ms/step - loss: 0.0166 - accuracy: 0.9946 - val_loss: 0.1161 -
val_accuracy: 0.9773
Epoch 6/12
1875/1875 [=====] - 3s 1ms/step - loss: 0.0151 - accuracy: 0.9948 - val_loss: 0.1241 -
val_accuracy: 0.9752
Epoch 7/12
1875/1875 [=====] - 3s 1ms/step - loss: 0.0121 - accuracy: 0.9959 - val_loss: 0.1189 -
val_accuracy: 0.9775
Epoch 8/12
1875/1875 [=====] - 3s 2ms/step - loss: 0.0152 - accuracy: 0.9952 - val_loss: 0.1121 -
val_accuracy: 0.9771
Epoch 9/12
1875/1875 [=====] - 3s 2ms/step - loss: 0.0117 - accuracy: 0.9965 - val_loss: 0.1323 -
val_accuracy: 0.9749
Epoch 10/12
1875/1875 [=====] - 3s 2ms/step - loss: 0.0115 - accuracy: 0.9961 - val_loss: 0.1200 -
val_accuracy: 0.9799
Epoch 11/12
1875/1875 [=====] - 3s 2ms/step - loss: 0.0100 - accuracy: 0.9969 - val_loss: 0.1282 -
val_accuracy: 0.9778
Epoch 12/12
1875/1875 [=====] - 3s 1ms/step - loss: 0.0124 - accuracy: 0.9959 - val_loss: 0.1238 -
val_accuracy: 0.9780
Out[33]: <keras.callbacks.History at 0x206dece8400>
```

Testing Model with custom image

```
In [35]: import cv2
```

```
In [60]: img = cv2.imread('5.png')
        plt.imshow(img)

Out[60]: <matplotlib.image.AxesImage at 0x206fc2abc10>
```



```
In [62]: img.shape

Out[62]: (237, 218, 3)

In [63]: img = cv2.resize((cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)), (28,28), interpolation=cv2.INTER_AREA)

In [64]: img.shape

Out[64]: (28, 28)

In [69]: import tensorflow

In [70]: img = tensorflow.keras.utils.normalize(img, axis=1)

In [71]: img = np.array(img).reshape(-1,28,28,1)

In [72]: img.shape

Out[72]: (1, 28, 28, 1)

In [74]: predictions = model.predict(img)
```

```
In [81]: print("Number Predicted:")
        print(np.argmax(model.predict(img)))

Number Predicted:
5
```

Model Accuracy

AS we can see from above model testing our model accuracy is 99%